

CS 536 Fall 2021

Lab 6: Application Layer Routing [220 pts]

Due: 12/1/2021 (Wed), 11:59 PM

Objective

In this lab, we will investigate basic techniques underlying application layer routing that are used to implement infrastructure network services such as virtual private networks (VPNs) and overlay networks. The key element is different forms of network tunneling that enable a measure of user control over packet forwarding and resultant routes that packets traverse. The creation of virtual networks over physical networks has myriad applications. For example, anonymity facilitated by overlay networks may be used by an hacker to hide its true location which slows down mitigation of an attack.

Reading

Read chapters 7, 8 and 9 from Peterson & Davie.

Problem 1 (220 pts)

1.1 Overall design

We will build an overlay network utilizing application layer routing to forward packets through paths that go through designated IP devices when packets are sent from source to destination. The executable binary, `zigzagrouter`, runs on a subset of Linux PCs in both lab machines (pod and amber). For example, `pod1-1.cs.purdue.edu`, `pod3-4.cs.purdue.edu`, `amber02.cs.purdue.edu`, `amber05.cs.purdue.edu`. For readability, the symbolic domain names of the IP addresses are used but your code should assume IP addresses in dotted decimal form. On each machine where `zigzagrouter` executes, three UDP sockets are created that bind to three unused ports. The port numbers are output to `stdout` so that the human operator knows which simplifies coding. At the sender, in our case the UDP ping client, `mydingcli`, from Problem 1 of lab2, an overlay network configuration binary, `zigzagconf`, is executed. To each IP device running `zigzagrouter`, `zigzagconf` sends a management UDP packet to the first port bound by the device informing it how to forward UDP packets arriving on the second and third ports bound by `zigzagrouter`. In networking parlance, we refer to the first port number as the control plane of the overlay network through which management information is communicated.

For example, `zigzagconf` running on `amber07.cs.purdue.edu` may send a UDP packet to `pod3-4.cs.purdue.edu` on the first port bound by `zigzagrouter` specifying

```
55000 39000 amber05.cs.purdue.edu 40000 50001 pod1-1.cs.purdue.edu
```

which instructs `zigzagrouter` on `pod3-4.cs.purdue.edu` to forward a UDP packet arriving on port 55000 (its second port) to `amber05.cs.purdue.edu:39000`. Similarly, UDP packets arriving on port 40000 (its third port) should be forwarded to `pod1-1.cs.purdue.edu` at port 50001. We refer to ports 55000 and 40000 at `pod3-4.cs.purdue.edu` as comprising part of the overlay network's data plane. The final destination of the `myding` UDP app from lab2 runs the `myding` server app, `mydingsrv`, say on `pod2-2.cs.purdue.edu` at port 55555. By sending instructions over the UDP-based control plane to each of the application layer routers `pod1-1.cs.purdue.edu`, `pod3-4.cs.purdue.edu`, `amber02.cs.purdue.edu`, and `amber05.cs.purdue.edu` they can be configured to forward `mydingcli` UDP packets from `amber07.cs.purdue.edu` to `mydingsrv` running on `pod2-`

2.cs.purdue.edu through four intermediate hops, say, amber02.cs.purdue.edu -> pod3-4.cs.purdue.edu -> amber05.cs.purdue.edu -> pod1-1.cs.purdue.edu.

The return path from mypingsrv running on pod2-2.cs.purdue.edu to mypingcli on amber07.cs.purdue.edu which uses the third port of the data plane need not be symmetric. For example, the return path may be pod1-1.cs.purdue.edu -> pod3-4.cs.purdue.edu -> amber05.cs.purdue.edu -> amber02.cs.purdue.edu. In this specific overlay network, the first hop (amber02.cs.purdue.edu) and last hop (pod1-1.cs.purdue.edu) of the forward data path will switch roles in the return data path. That is, pod1-1.cs.purdue.edu is the first hop of the return path and amber02.cs.purdue.edu the last hop of the return path. Since mypingcli binds to an ephemeral port number and reuse mypingcli as a legacy app as is (without changing its code to output its port to stdout), zigzagconf will send the following UDP management packet to the first hop on the forward path amber02.cs.purdue.edu

```
25000 55000 pod3-4.cs.purdue.edu 26000 0 0.0.0.0
```

In the above, 25000 represents the second port bound by amber02.cs.purdue.edu and 26000 its third port. The values 0 and 0.0.0.0 specify that zigzagrouter running on amber02.cs.purdue.edu should remember the IPv4 address and port number of the first UDP packet received on port 25000 (i.e., forward path) and replace 0.0.0.0 and 0 with their values so that the response sent from mypingsrv can be delivered to mypingcli running on amber07.cs.purdue.edu. To forward a UDP means to extract its payload and transmit it as the payload of a new UDP from application layer router to its next hop (an application layer router, source, or final destination).

1.2 Legacy compatibility

The UDP ping server, mypingsrv, runs as is, and the UDP ping request will appear as originating from the last hop pod1-1.cs.purdue.edu in the above example. The UDP ping client, mypingcli, does not run fully in a legacy compatible manner since its command-line argument must specify the coordinates of the first hop in the forward path. That is,

```
% mypingcli amber07.cs.purdue.edu amber02.cs.purdue.edu 25000
```

instead of the final destination's coordinates. In production systems, functionality of zigzagrouter are implemented as dynamically loadable kernel modules in operating systems such as Linux and Windows. The kernel modules act as hooks in its protocol stack that allow intercepting and modifying (or mangling) messages depending on where the hooks are installed. For example, in a simple VPN where a single intermediate hop acts as the conduit -- i.e., packets are not bounced around multiple times as in a pinball machine -- a legacy client app such as a web browser works as is with final destination IP addresses. However, transparent to the client app the kernel module will tunnel client messages through the VPN server which will then appear to the server that the request is originating from the VPN server.

In the case of Purdue's campus network, some services are only honored if the source IP address belongs to one of the Purdue's CIDR IP prefixes. By default, these services would not be accessible off-campus through ISPs since their IP addresses are not part of Purdue's IP prefix blocks. An on-campus VPN server that acts as a conduit can allow the same services to be accessed off-campus. Since modifying kernel behavior to implement network protocols is outside the scope of the lab assignments, we are approaching the problem from a less transparent and elegant manner. Linux and Windows kernel programming to implement these features is not difficult but requires additional background and skills beyond system programming.

1.3 Implementation details

The overlay network configuration app, zigzagconf, should read from a configuration file (ASCII text file), zzoverlay.dat, to whom on the control plane what forwarding configuration information should be transmitted. For example, one entry in zzovrelay.dat for the application layer router zigzagrouter running on pod3-4.cs.purdue.edu would be

```
pod3-4.cs.purdue.edu 33333
```

```
55000 39000 amber05.cs.purdue.edu
40000 50001 pod1-1.cs.purdue.edu
```

where pod3-4.cs.purdue.edu specifies the IP address (in dotted decimal form) of an application layer router and its first (i.e., control plane) port number 33333, and 55000 39000 amber05.cs.purdue.edu 40000 50001 pod1-1.cs.purdue.edu specifying the forwarding instructions on the forward and return paths. The first entry of zzoverlay.dat is an integer specifying the number of application layer routers, followed by three lines separated by newline. The IP addresses (dotted decimal) and port numbers should be separated by spaces.

A UDP management packet transmitted over the control plane has the format: 2 bytes for second port number, 2 bytes for next hop port number, 4 bytes for IPv4 address for the forward path, 2 bytes for third port number, 2 bytes for next hop port number, 4 bytes for IPv4 address for the return path. Hence a total of 16 bytes.

Whenever a control plane message is received, zigzagrouter should output to stdout a timestamp followed by the content of the message after installing it in a data structure that represents its forwarding table. When a data plane packet is received, zigzagrouter should output to stdout a timestamp followed the sender's IP address (in dotted decimal form) and port number, and the payload of the UDP packet. When zigzagconf transmits a control plane UDP packet, print to stdout the timestamp, the destination IP address and port number, and the payload of the packet. The output will help monitor and debug zigzagrouter and zigzagconf actions. When zigzagrouter executes, it creates a socket to be used for the control plane. Instead of binding to an ephemeral port number, provide the port number to be used for the control plane as a command-line argument of zigzagrouter.

zigzagrouter needs to handle three socket descriptors (1 for control plane, 2 for data plane). Use the select() system call which facilitates monitoring events on multiple descriptors.

1.4 Testing

Test the overlay network application that utilizes multiple hops across both pod and amber lab machines. Make sure to use IPv4 addresses in dotted decimal form, not their symbolic (i.e., domain) names. That is, 128.10.25.214 in place of pod3-4.cs.purdue.edu. Run the UDP myping app without intermediate application layer routers to gauge RTT. Then run the app with the same parameters specified in pingparam.dat (keep it simple). Test and verify that the overlay network application works correctly. The UDP ping app should report larger RTT values that increase with the number of hops, but otherwise work as before. Discuss your finding in lab6.pdf. Submit your work in v1/ following the convention of previous labs (e.g., Makefile, README).

Note: This problem may be tackled as a group effort involving up to 3 people. If going the group effort route, please specify in lab6.pdf on its cover page who the members are, who did what work including programming the various components, performance evaluation, and write-up. If you participated in a group effort in lab5, the members of lab6 cannot overlap. Whether you implement lab6 as an individual effort or make it a group effort is up to you. Keep in mind the trade-offs: group effort incurs coordination overhead which can slow down execution, especially for a 2-week assignment. Benefits include collaborative problem solving and some parallel speed-up if efficiently executed. Regarding late days, for a group to use k ($= 1, 2, 3$) late days, every member of the group must have k late days left.

Note: For students who have earned 100+ extra credits (bonus problems in lab1-lab5 and late days unused), you may consider solving a simpler problem where instead of zigzagconf communicating forwarding instructions over the control plane, the information is read from a configuration file, zzone.dat, and support for a single application layer router (i.e., simple VPN) is sufficient. For example, if pod3-4.cs.purdue.edu were the application layer router where zigzagrouter runs, zzone.dat would contain

```
55000 39000 128.10.112.135
40000 0 0.0.0.0
```

Thus when zigzagrouter starts it reads the forwarding information from zzone.dat and is ready to forward UDP packets arriving on the data plane. The simplified problem counts as 120 points toward lab6's total of

220 points. Hence solving the simplified problem plus 100 points of extra credit earned would yield 220 points. If you decide to solve the simplified problem, please indicate so on the cover page of lab6.pdf. In a group effort, all members must solve the same problem, full or simplified.

Turn-in Instructions

Electronic turn-in instructions:

We will use turnin to manage lab assignment submissions. Go to the parent directory of the directory lab6/ where you deposited the submissions and type the command

```
turnin -v -c cs536 -p lab6 lab6
```

This lab is an individual effort. Please note the assignment submission policy specified on the course home page.

[Back to the CS 536 web page](#)